

LSM Tree Implementation

Nicolas Drizard

TF: Lukas Maas
May, 3rd 2015

Outline

- 1 Introduction
- 2 Design
 - Data Structure
 - Operations
 - Persistency
 - Implementation details
 - Parallel Reads
 - Alternative Merging Strategy
- 3 Experiments
- 4 Future Steps

Table of Contents

- 1 Introduction
- 2 Design
 - Data Structure
 - Operations
 - Persistency
 - Implementation details
 - Parallel Reads
 - Alternative Merging Strategy
- 3 Experiments
- 4 Future Steps

LSM-Tree




Figure: LSM Tree structure

Table of Contents

- 1 Introduction
- 2 Design
 - Data Structure
 - Operations
 - Persistency
 - Implementation details
 - Parallel Reads
 - Alternative Merging Strategy
- 3 Experiments
- 4 Future Steps

Layers Architecture

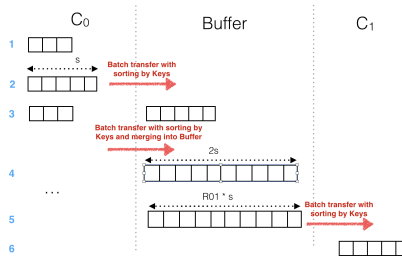


Figure: Merging Strategy

Operations

- READ
 - ① Linear search in C_0
 - ② binary search in other components
- APPEND (insert, update, delete)
 - ① Linear search in C_0 (removing/updating keys)
 - ② Merging if full

Each component contains two arrays: keys and values.
 C_0 and buffer on memory

Optimizations

- Memory Mapped File
- Bounds Check
- Fault tolerance strategy
- Bloom Filter

Parallel Design

Two approaches are considered:

- Killing threads operating on deeper component once the key is found in its lowest component.
- Threads communication through a shared variable storing the lowest component on which the key has been found so far




Figure: Alternative Merging Strategy

Table of Contents

- 1 Introduction
- 2 Design
 - Data Structure
 - Operations
 - Persistency
 - Implementation details
 - Parallel Reads
 - Alternative Merging Strategy
- 3 Experiments
- 4 Future Steps

Set Up

- 1 000 000 keys inserted
- keys are int, values string of size 32

Set Up

5 different LSM structures used:

- ① C_0 size = 1000; size ratios = 3
- ② C_0 size = 2000; size ratios = 3
- ③ C_0 size = 1000; size ratios = 3 and then 9
- ④ C_0 size = 1000; size ratios = 5
- ⑤ C_0 size = 500; size ratios = 5

Results




Figure: Uniform Reads




Figure: Uniform Updates

Results




Figure: Skewed reads (20% start)




Figure: Skewed updates (20% start)

Results




Figure: Skewed reads (20% end)




Figure: Skewed updates (20% end)

Results

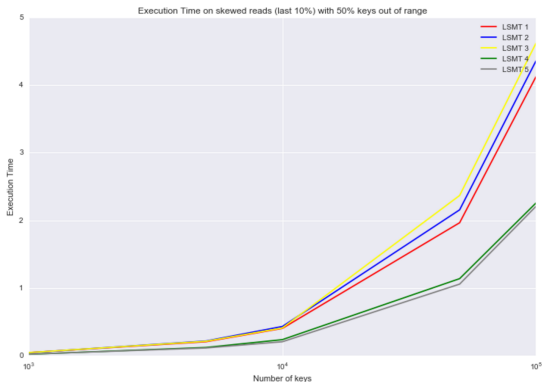


Figure: Skewed reads (10% end) and 50% out of range

Table of Contents

- 1 Introduction
- 2 Design
 - Data Structure
 - Operations
 - Persistency
 - Implementation details
 - Parallel Reads
 - Alternative Merging Strategy
- 3 Experiments
- 4 Future Steps

- Experiments on Bloom Filter and Parallel implementation
- Possibility to choose the merging strategy